

Score Based Models

Advanced Topics in Deep Learning

Vahid Tarokh

ECE 590-02

Spring 2022

Partition Function

- In the previous lecture, we discussed energy-based models.
- The main issue with these models was the calculation involving the partition function.
- We considered Monte Carlo approximation methods both for the training and testing of energy models and in particular those based on deep neural networks.
- The results were quite good. We wonder how better it can be if we did not have to approximate.
- In this lecture, we will discuss Fisher Score based approaches to address the partition function.
- First we discuss general score functions.

Score Functions

- How to evaluate the cost/utility of models fit to data?
 - Probability density function p and data observations y
 - Score function $s: (y, p) \mapsto s(y, p)$
 - The smaller, the better
 - Example: logarithmic scoring function $s(y, p) = -\log p(y)$
- Score rule
 - Logarithmic rule
 - Minimize $s(y, p) = -\log p(y)$ (here y can be a vector)
 - MLE
 - $p \equiv p_\theta$, obtain the θ that achieves the minimum of $s(y, p)$
 - Widely used

Overview of Logarithmic Score

- Logarithmic Score

- Widely adopted for statistical inference, information theory, etc.
- Examples: maximum likelihood estimation, hypothesis test, Bayesian model comparison

- Why is it so popular?

- Elementary function of the probability density function
- Intimate relation with KL-divergence

- Minimizing

- $n^{-1} \sum_i s(y_i, p) = n^{-1} \sum_i \{-\log p(y_i)\}$

- is asymptotically equivalent to minimizing

- $E_*\{-\log p(y)\} = D_{KL}(p_*||p)$

- (under i.i.d. or ergodicity assumptions)

- In a well-specified scenario: optimum $p = p_*$

- In a mis-specified scenario: optimum p is the closest to p_* in KL sense (under some regularity conditions)

p_* : true distribution

E_* : expectation w.r.t. p_*

Maximum Likelihood Estimator (MLE)

Notice that

$$-\frac{\sum_1^n \log p_\theta(y_i)}{n} \rightarrow E_{p^*} [-\log p_\theta(y)]$$

- Minimizing

$$-\frac{\sum_1^n \log p_\theta(y_i)}{n}$$

is asymptotically equivalent to minimizing

$$E_*\{-\log p_\theta(y)\} = D_{KL}(p_*||p_\theta) + H(p_*)$$

or equivalently

$$D_{KL}(p_*||p_\theta).$$

Some Thoughts

- Up to now, we mainly used MLE, cross-entropy, etc. for the training of our models.
- The motivation as to minimize KL divergence between the true model and the postulated modeled (e.g. DNN based).
- This was equivalent to minimizing the logarithmic score (as discussed)
- By doing this, we ran into the curse of partition function.
 - We did Monte-Carlo approximation to address this.
- Not clear that logarithmic score generates the best images perceptually.
- May be we can use another score that does not have the curse of partition function?
- If possible, then we must minimize another distance with the true model

Fisher Divergence and the Fisher Score

- Fisher score from Fisher divergence

- Consider minimizing the following **Fisher** divergence instead of KL

$$\frac{1}{2} E_* \|\nabla_y \log p(y) - \nabla_y \log p_*(y)\|^2 \quad (1)$$
$$= E_* \{s_P(y, p)\} + c_*$$

$$s_P(y, p) = \frac{1}{2} \|\nabla_y \log p(y)\|^2 + \Delta_y \log p(y),$$

- c_* only depends on p_* , the true data-generating distribution. (prove this)
- The minimum, zero, is achieved if and only if $p(y) = p_{\theta_*}(y)$
 p_{θ_*} : the closest density to p_* in the sense of (1)
 θ_* : true parameter (if well-specified)
- Suppose we minimize $n^{-1} \sum_i s_P(y_i, p_\theta)$ **(e.g. using SGD)**
 - It is approximately minimizing $E_* \{s_P(y, p_\theta)\}$
 - It leads to $\theta \rightarrow \theta_*$ in probability



Ronald Aylmer Fisher

Desirable properties of Fisher-Score

- A **simple** function of $p(y)$ and its derivatives
- **Proper** in the sense that $E_*\{s(y, p)\}$ is only minimized at $p = p_*$ a.s. under mild conditions
- **Scale-invariance** in the sense that $E_*\{s(y, p)\} \equiv E_*\{s(y, c p)\}$
 p does not need to be a density function as long as c does not depend on y
- This means that calculations can be done with unnormalized distributions. **Partition function will not be playing a role at all!**
 - This is unlike the logarithmic score.
- May be defined for **improper priors**.
 - The energy model does not even need to be integrable!

Relationship Between Fisher and KL Divergences.

- Relationship between Fisher Divergence and KL

$$D_{\nabla}(p, q) = -2 \frac{d}{dt} D_{\text{KL}}(p_t, q_t) \big|_{t=0}$$

- $X \sim p(\cdot)$
- $Y \sim q(\cdot)$
- $X_t = X + \mathcal{N}(0, t)$
- $Y_t = Y + \mathcal{N}(0, t)$ *perturbed by independent $\mathcal{N}(0, t)$*
- $X_t \sim p_t(\cdot)$
- $Y_t \sim q_t(\cdot)$
- The proof is similar (and follows from) de Bruijn equality. $J_F(X) = 2 \frac{d}{dt} [H(X_t)]$

Incorporation into AI systems

- Fisher score can be used for
 - Parameter inference
 - Sequential Monte Carlo
 - Bayesian model comparison
 - Causality Calculations
 - Change detection
 - **Training Energy Models (Topic of current lecture)**
- Fisher score may have some advantages over logarithmic scores in some cases
 - both computational and foundational.

Energy models

- We can think of the the function given by NN_{θ} as $g_{\theta}(y)$ and the corresponding value of partition as Z_{θ} . Thus

$$p_{\theta}(y) = \frac{e^{-g_{\theta}(y)}}{Z_{\theta}}.$$

- Calculating the Fisher score:

$$\nabla_y \ln p(y) = -\nabla_y g(y)$$

$$\Delta_y \ln p(y) = -\Delta_y g(y).$$

- The Fisher score is given by

$$\frac{1}{2} \|\nabla_y g(y)\|^2 - \Delta_y g(y)$$

- This does not require the calculation of the partition function.
- We will present some examples.

Gauss-Bernoulli Boltzmann machine


- Continuous data modeled by Restricted Boltzmann machine (RBM) (see ECE685D Notes)
- An RBM consists of m visible **real valued** units \mathbf{V} and n binary hidden units \mathbf{H}
- The energy function is

$$E_{\theta}(V, H) = -H^T W V - C^T V - B^T H + \frac{1}{2} V^T V$$

- The probability density is

$$p_{\theta}(V, H) = \exp(H^T W V + C^T V + B^T H - \frac{1}{2} V^T V)$$

- $\theta = (\mathbf{W}, \mathbf{C}, \mathbf{B})$ denote all the parameters of the model



Gauss-Bernoulli Boltzmann Machine

- The Fisher score function can be easily calculated.

$$s_{\theta}(V) = \frac{1}{2} ||W^T \rho + C + V||^2 + ||\text{diag}(\rho')W||^2$$

where $\rho = \sigma(WV + B)$, $\rho' = \sigma'(WV + B)$

- We would like to minimize the sum of Fisher score over all training data points. To this end, we use **stochastic gradient descent**.
- This means that we need to calculate the derivatives of score function with respect to θ .
- This calculation is straightforward but somewhat painful. The results can be put in closed form and involve the sigmoid functions and its derivatives.
 - Exercise: Verify these calculations.
- We use these derivatives for the training with SGD.


Training of Restricted Boltzmann Machine

- Continuous data modeled by Restricted Boltzmann machine (RBM)
 - An RBM consists of 784 visible units x and 30 hidden units h


At the same computational complexity, our approach creates a generative model that is closer to the underlying truth



Original Images



Generated Images
using Fisher



Generated Images
using Classical KL
Divergence
(approximated by
Contrastive Divergence)

Video: Training of RBMs using New versus Classical Approaches

Special Energy Models

- Some special energy models are particularly suitable for the application of Fisher score.

- Consider **the exponential family** (see ECE685D notes)

$$p_{\theta}(y) = \exp\{a(y) + b(\theta) + \theta^T s(y)\}$$

- By minimizing the Fisher score $\sum_{t=1}^n s_P(y_t, p_{\theta})$
- we can have the **closed form solution** (usually not possible with MLE)

$$\hat{\theta}_H = - \left(\sum_{t=1}^n J(y_t)^T J(y_t) \right)^{-1} \left\{ \sum_{t=1}^n \Delta s(y_t) + J(y_t)^T \nabla a(y_t) \right\}$$

where $J(y)$ denotes the matrix $J(y) = \left[\frac{\partial s_j(y)}{\partial y_i} \right]_{i,j}$ and $\Delta s(y)$ denotes the vector $[\Delta s_j(y)]_j$

- No need to apply numerical techniques.

Deep Neural Networks as Energy Models

- We can think of the the function given by NN_{θ} as $g_{\theta}(y)$.

- To calculate the Fisher Score, we need to calculate

$$\frac{1}{2} |\nabla_y g(y)|^2 - \Delta_y g(y)$$

- The first part is easy to calculate, but the Laplacian may be computationally complex to calculate.
- Thus it is not immediately clear how to apply the Fisher score for the training of energy models where the energy function is given by a large deep neural network.
- We must think of some other remedies.

Key idea

- Remember that we started from the Fisher divergence with data the true pdf of data $p(x)$:

$$\frac{1}{2} E_* \|\nabla_x \log p_\theta(x) - \nabla_x \log p(x)\|^2 = E_* \{s_P(x, p_\theta)\} + c$$

- In a sense all we have been trying to do is to come up with a model $p_\theta(x)$ whose $\nabla_x \log p_\theta(x)$ is close to $\nabla_x \log p(x)$.
 - This led to Fisher score.
- If we estimate $\nabla_x \log p(x)$ well, then we can use MALA to create new images that look like the original images.
 - The estimate $\nabla_x \log p(x)$ is all MALA needs.
- Now suppose we made a deep neural network NN_θ whose output $s_\theta(x)$ estimates $\nabla_x \log p(x)$, then we are in business!
- This idea applies beyond energy models.**
- It remains to discuss how to construct such a deep neural network whose output

$$s_\theta(\mathbf{x}) \approx \nabla_{\mathbf{x}} \log p(\mathbf{x})$$

Sliced Score Matching

- An approach is sliced score matching [Song et al. UAI 2019].
- The **slicing** idea is to project a high dimensional optimization objective into random lower dimensional subspaces (in most cases only one dimensional subspaces) and optimize these sums.
- To this end, we generate some unit vectors \mathbf{v} according to probability $p_{\mathbf{v}}(\mathbf{v})$. In most case, we want this distribution to be uniform over all unit directions. This can be achieved by drawing samples from $N(0, I)$ and normalizing the samples to have unit length.

- Now we calculate the sliced Fisher divergence:

$$\frac{1}{2} E_{\mathbf{v} \sim p_{\mathbf{v}}} E_{\mathbf{x} \sim p_{\text{data}}} [(\mathbf{v}^T \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) - \mathbf{v}^T \mathbf{s}_{\theta}(\mathbf{x}))^2]$$

- This is the **sliced** surrogate of the Fisher divergence that we originally optimized.

Sliced Fisher Divergence

- Just as in Fisher divergence, we expand and integrate by parts and observe that this is equivalent to minimizing

$$E_{\mathbf{v} \sim p_{\mathbf{v}}} E_{\mathbf{x} \sim p_{\text{data}}} \left[\mathbf{v}^{\top} \nabla_{\mathbf{x}} s_{\theta}(\mathbf{x}) \mathbf{v} + \frac{1}{2} (\mathbf{v}^{\top} s_{\theta}(\mathbf{x}))^2 \right]$$

- We can calculate this empirically by generating various directions (unit vectors) \mathbf{v} according to probability $p_{\mathbf{v}}(\mathbf{v})$ at each step of SGD and for the data point calculate the empirical average. Then we can do descent on θ .
- If $s_{\theta}(x)$ is the output of a deep neural network $NN_{\theta}(x)$ then calculating $\nabla_x s_{\theta}(x)$ is doable (in fact Pytorch does it).
- The whole process is summarized in the algorithm given in the next page.

Sliced Score Matching

Sample a minibatch of datapoints $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\} \sim p_{\text{data}}(\mathbf{x})$

Sample a minibatch of projection directions $\{\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n\} \sim p_{\mathbf{v}}$

Estimate the sliced score matching loss with empirical means

$$\frac{1}{n} \sum_{i=1}^n \left[\mathbf{v}_i^T \nabla_{\mathbf{x}} \mathbf{s}_{\theta}(\mathbf{x}_i) \mathbf{v}_i + \frac{1}{2} (\mathbf{v}_i^T \mathbf{s}_{\theta}(\mathbf{x}_i))^2 \right]$$

- Note: We can use more than one projections per datapoint too.
- Use Stochastic Gradient Descent to minimize the cost

Denoising Score Matching

- Denoising score matching [Vincent 2011] is another way to train DNNs based on matching the score of a noise-perturbed distribution of the data generating distribution $p_{\text{data}}(\mathbf{x})$
- Noise is added to training images x to arrive at the noisy images \tilde{x} .

$$p_{\text{data}}(\mathbf{x}) \longrightarrow q_{\sigma}(\tilde{\mathbf{x}} \mid \mathbf{x}) \longrightarrow q_{\sigma}(\tilde{\mathbf{x}})$$

- Typically the noise is Gaussian $N(0, \sigma^2)$.
- Now we try to have $s_{\theta}(x)$ (the output of DNN) match $q_{\sigma}(\tilde{\mathbf{x}})$
- We do some calculations next.

Some Calculations

- We have

$$\frac{1}{2} E_{\tilde{\mathbf{x}} \sim q_{\sigma}} [\|\nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}}) - \mathbf{s}_{\theta}(\tilde{\mathbf{x}})\|_2^2] = \frac{1}{2} \int q_{\sigma}(\tilde{\mathbf{x}}) \|\nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}})\|_2^2 d\tilde{\mathbf{x}} + \frac{1}{2} \int q_{\sigma}(\tilde{\mathbf{x}}) \|\mathbf{s}_{\theta}(\tilde{\mathbf{x}})\|_2^2 d\tilde{\mathbf{x}} \\ - \int q_{\sigma}(\tilde{\mathbf{x}}) \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}})^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}}$$

- The first blue term is independent of training of the DNN:

$$= \text{const.} + \frac{1}{2} E_{\tilde{\mathbf{x}} \sim q_{\sigma}} [\|\mathbf{s}_{\theta}(\tilde{\mathbf{x}})\|_2^2] - \int q_{\sigma}(\tilde{\mathbf{x}}) \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}})^{\top} \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}}$$

- Next we manipulate the second (red) term.

$$\begin{aligned}
& - \int q_\sigma(\tilde{\mathbf{x}}) \nabla_{\tilde{\mathbf{x}}} \log q_\sigma(\tilde{\mathbf{x}})^\top \mathbf{s}_\theta(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
&= - \int q_\sigma(\tilde{\mathbf{x}}) \frac{1}{q_\sigma(\tilde{\mathbf{x}})} \nabla_{\tilde{\mathbf{x}}} q_\sigma(\tilde{\mathbf{x}})^\top \mathbf{s}_\theta(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
&= - \int \nabla_{\tilde{\mathbf{x}}} q_\sigma(\tilde{\mathbf{x}})^\top \mathbf{s}_\theta(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
&= - \int \nabla_{\tilde{\mathbf{x}}} \left(\int p_{\text{data}}(\mathbf{x}) q_\sigma(\tilde{\mathbf{x}} | \mathbf{x}) d\mathbf{x} \right)^\top \mathbf{s}_\theta(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
&= - \int \left(\int p_{\text{data}}(\mathbf{x}) \nabla_{\tilde{\mathbf{x}}} q_\sigma(\tilde{\mathbf{x}} | \mathbf{x}) d\mathbf{x} \right)^\top \mathbf{s}_\theta(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
&= - \int \left(\int p_{\text{data}}(\mathbf{x}) q_\sigma(\tilde{\mathbf{x}} | \mathbf{x}) \nabla_{\tilde{\mathbf{x}}} \log q_\sigma(\tilde{\mathbf{x}} | \mathbf{x}) d\mathbf{x} \right)^\top \mathbf{s}_\theta(\tilde{\mathbf{x}}) d\tilde{\mathbf{x}} \\
&= - \iint p_{\text{data}}(\mathbf{x}) q_\sigma(\tilde{\mathbf{x}} | \mathbf{x}) \nabla_{\tilde{\mathbf{x}}} \log q_\sigma(\tilde{\mathbf{x}} | \mathbf{x})^\top \mathbf{s}_\theta(\tilde{\mathbf{x}}) d\mathbf{x} d\tilde{\mathbf{x}} \\
&= - E_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x}), \tilde{\mathbf{x}} \sim q_\sigma(\tilde{\mathbf{x}} | \mathbf{x})} [\nabla_{\tilde{\mathbf{x}}} \log q_\sigma(\tilde{\mathbf{x}} | \mathbf{x})^\top \mathbf{s}_\theta(\tilde{\mathbf{x}})]
\end{aligned}$$

Computation Continued

- This in turn equals to

$$\begin{aligned} &= \text{const.} + \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x}), \tilde{\mathbf{x}} \sim q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})} [\|\mathbf{s}_{\theta}(\tilde{\mathbf{x}}) - \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})\|_2^2] \\ &\quad - \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x}), \tilde{\mathbf{x}} \sim q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})} [\|\nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})\|_2^2] \end{aligned}$$

which is

$$\begin{aligned} &= \text{const.} + \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x}), \tilde{\mathbf{x}} \sim q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})} [\|\mathbf{s}_{\theta}(\tilde{\mathbf{x}}) - \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})\|_2^2] + \text{const.} \\ &= \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x}), \tilde{\mathbf{x}} \sim q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})} [\|\mathbf{s}_{\theta}(\tilde{\mathbf{x}}) - \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x})\|_2^2] + \text{const.} \end{aligned}$$

Denoising Score Matching

- Thus we proved that

$$\begin{aligned} & \frac{1}{2} E_{\tilde{\mathbf{x}} \sim p_{\text{data}}} [\| \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) - \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}}) \|_2^2] \\ &= \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x}), \tilde{\mathbf{x}} \sim q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})} [\| \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) - \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}} | \mathbf{x}) \|_2^2] + \text{const.} \end{aligned}$$

- Now we observe that $\nabla_{\tilde{x}} \log q_{\sigma}(\tilde{x}|x)$ is easy to calculate. In fact assuming Gaussian noise $N(0, \sigma^2)$ this is equal to $\frac{-(\tilde{x} - x)}{\sigma^2}$.
- This means that the above distance can be computed empirically even in very high dimensions using backpropagation.
- The caveat is that we are matching the $\nabla_{\tilde{x}} \log q_{\sigma}(\tilde{x}|x)$ and not $\nabla_x \log p(x)$.
- Let us summarize this algorithm [Vincent 2011] in the next slide.

Denoising Score Matching (pseudocode)

- Choose a Gaussian $N(0, \sigma^2)$ or any other noise $q_\sigma(\tilde{x}|x)$ with small variance σ and conditional mean x .

Sample a minibatch of datapoints $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\} \sim p_{\text{data}}(\mathbf{X})$

Sample a minibatch of perturbed datapoints $\{\tilde{\mathbf{x}}_1, \tilde{\mathbf{x}}_2, \dots, \tilde{\mathbf{x}}_n\} \sim q_\sigma(\tilde{\mathbf{X}})$

$$\tilde{\mathbf{x}}_i \sim q_\sigma(\tilde{\mathbf{x}}_i | \mathbf{x}_i)$$

Estimate the denoising score matching loss with empirical means

$$\frac{1}{2n} \sum_{i=1}^n [\| \mathbf{s}_\theta(\tilde{\mathbf{x}}_i) - \nabla_{\tilde{\mathbf{x}}} \log q_\sigma(\tilde{\mathbf{x}}_i | \mathbf{x}_i) \|_2^2]$$

If Gaussian perturbation

$$\frac{1}{2n} \sum_{i=1}^n \left[\left\| \mathbf{s}_\theta(\tilde{\mathbf{x}}_i) + \frac{\tilde{\mathbf{x}}_i - \mathbf{x}_i}{\sigma^2} \right\|_2^2 \right]$$

- Use Stochastic Gradient Decent (SGD) to minimize over θ .

Fine-tuning σ

- Let us calculate the effect of σ on the loss function using reparameterization trick (please see ECE685D Lecture notes).

$$\begin{aligned} \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}} E_{\tilde{\mathbf{x}} \sim q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})} \left[\left\| \mathbf{s}_{\theta}(\tilde{\mathbf{x}}) + \frac{\tilde{\mathbf{x}} - \mathbf{x}}{\sigma^2} \right\|_2^2 \right] &= \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}} E_{\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})} \left[\left\| \mathbf{s}_{\theta}(\mathbf{x} + \sigma \mathbf{z}) + \frac{\mathbf{z}}{\sigma} \right\|_2^2 \right] \\ &= \frac{1}{2} E_{\mathbf{x} \sim p_{\text{data}}} E_{\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})} \left[\left\| \mathbf{s}_{\theta}(\mathbf{x} + \sigma \mathbf{z}) \right\|_2^2 + 2 \mathbf{s}_{\theta}(\mathbf{x} + \sigma \mathbf{z})^{\top} \frac{\mathbf{z}}{\sigma} + \frac{\left\| \mathbf{z} \right\|_2^2}{\sigma^2} \right] \end{aligned}$$

- Now if $\sigma \rightarrow 0$ we can easily see that the variances of $\frac{\mathbf{z}}{\sigma}$ and $\frac{\left\| \mathbf{z} \right\|_2^2}{\sigma^2}$ both goes to ∞ .
- So we can not let $\sigma \rightarrow 0$ as we will encounter computational instability in training.
 - It is important to fine-tune σ .

Fine-tuning σ

- One way to do this [Song and Ermon 2019] is to choose a decreasing sequence of positive numbers $\sigma_1 > \sigma_2 > \dots > \sigma_L$.
- Then use the Denoising Score Matching (or the sliced score matching algorithm applied to noisy data) with Gaussian noise variance $N(0, \sigma_i^2)$ to train score function $s_\theta(x, \sigma_i)$ at i -th step.
- Then run MALA with appropriately selected noise variance to sample.
- Use the samples to train for the next $(i + 1)$ -th step.
- Details are given in the next page.

Annealed Score Matching and MALA

```
1: Initialize  $\tilde{\mathbf{x}}_0$ 
2: for  $i \leftarrow 1$  to  $L$  do
3:    $\alpha_i \leftarrow \epsilon \cdot \sigma_i^2 / \sigma_L^2$   $\triangleright \alpha_i$  is the step size.
4:   for  $t \leftarrow 1$  to  $T$  do
5:     Draw  $\mathbf{z}_t \sim \mathcal{N}(0, I)$ 
6:      $\tilde{\mathbf{x}}_t \leftarrow \tilde{\mathbf{x}}_{t-1} + \frac{\alpha_i}{2} \mathbf{s}_\theta(\tilde{\mathbf{x}}_{t-1}, \sigma_i) + \sqrt{\alpha_i} \mathbf{z}_t$ 
7:   end for
8:    $\tilde{\mathbf{x}}_0 \leftarrow \tilde{\mathbf{x}}_T$ 
9: end for
   return  $\tilde{\mathbf{x}}_T$ 
```

Selection of sequence of noise variances

- σ_1 can be selected as the maximum Euclidean distance taken over pairs of data points.
- It has been heuristically observed that $\sigma_1 > \sigma_2 > \dots > \sigma_L$ can be selected to form a geometrically decaying sequence with the decay rate and L to be fine-tuned.
- Yet another heuristic method is to minimize the weight loss function (recall the reparameterization trick)

$$\frac{1}{L} \sum_{i=1}^L \lambda(\sigma_i) E_{\mathbf{x} \sim p_{\text{data}}, \mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})} \left[\left\| \mathbf{s}_{\theta}(\mathbf{x} + \sigma_i \mathbf{z}, \sigma_i) + \frac{\mathbf{z}}{\sigma_i} \right\|_2^2 \right]$$

where the weight $\lambda(\sigma_i) = \sigma_i^2$

- Combining all these heuristics, we arrive at the following algorithm [Song and Ermon 2019].

Noise conditional score networks Training Algorithm

Sample a mini-batch of datapoints $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\} \sim p_{\text{data}}$

Sample a mini-batch of noise scale indices

$$\{i_1, i_2, \dots, i_n\} \sim \mathcal{U}\{1, 2, \dots, L\}$$

Sample a mini-batch of Gaussian noise $\{\mathbf{z}_1, \mathbf{z}_2, \dots, \mathbf{z}_n\} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

Estimate the weighted mixture of score matching losses

$$\frac{1}{n} \sum_{k=1}^n \left[\|\sigma_{i_k} \mathbf{s}_{\theta}(\mathbf{x}_k + \sigma_{i_k} \mathbf{z}_k, \sigma_{i_k}) + \mathbf{z}_k\|_2^2 \right]$$

- Use Stochastic Gradient descent to minimize.

Experiments [Song and Ermon 2019]

Model	Inception	FID
CIFAR-10 Unconditional		
PixelCNN [59]	4.60	65.93
PixelIQN [42]	5.29	49.46
EBM [12]	6.02	40.58
WGAN-GP [18]	$7.86 \pm .07$	36.4
MoLM [45]	$7.90 \pm .10$	18.9
SNGAN [36]	$8.22 \pm .05$	21.7
ProgressiveGAN [25]	$8.80 \pm .05$	-
NCSN (Ours)	$8.87 \pm .12$	25.32

Experiments: High Resolution Image Generation [Song and Ermon 2019]

